

LLM-Powered Confluence Search and Q&A Tool — Technical Design

Version: 1.0.0 | **Date:** February 26, 2026 | **Owner:** AI Engineering - Lorant Zsarnowszky

Executive Summary

This document outlines a production-grade architecture for an LLM-powered Confluence search and Q&A system using Retrieval-Augmented Generation (RAG). The solution addresses four core challenges: maintaining up-to-date knowledge synchronization, handling multimodal attachments (images, diagrams, documents), managing multi-space context with project-specific metadata, and continuously evaluating answer quality.

Key Technical Pillars:

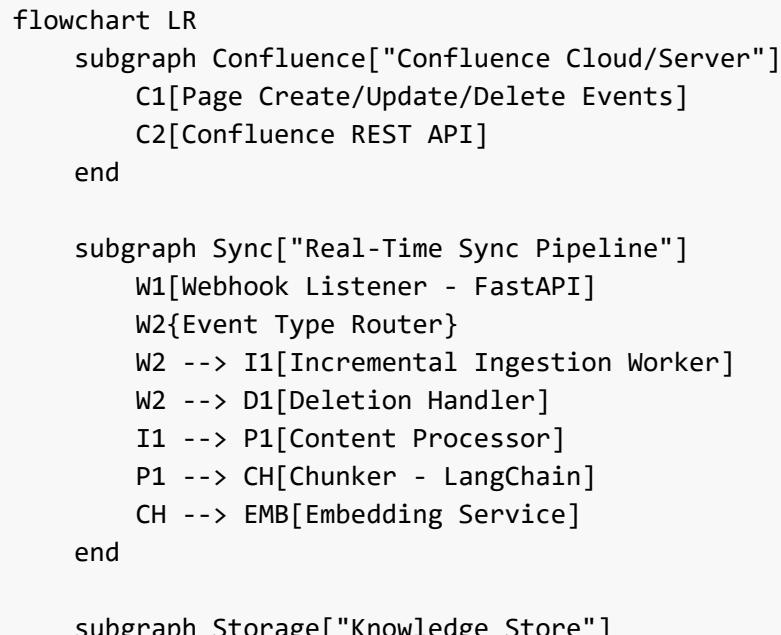
- **Real-time sync** via webhook-driven incremental ingestion pipeline
- **Multimodal RAG** supporting text, images (via vision LLMs), and document attachments
- **Space-aware retrieval** with metadata filtering and project-specific context injection
- **Automated evaluation** using LLM-as-judge + human feedback loops

1. Keeping Confluence Knowledge Accessible and Up-to-Date

1.1 Challenge

Confluence content evolves continuously — pages are created, updated, archived, and deleted. The RAG system must reflect the current state without manual re-indexing or stale results.

1.2 Solution Architecture



```

V1[(Qdrant Vector DB
with HNSW index)]
M1[(PostgreSQL
Metadata + Full Text)]
end

C1 --> W1
C2 --> I1
EMB --> V1
P1 --> M1
D1 --> V1
D1 --> M1

```

1.3 Implementation Strategy

Real-Time Webhook Integration

- **Confluence Webhooks:** Register webhooks for `page_created`, `page_updated`, `page_trashed`, `page_removed` events
- **Webhook Listener:** FastAPI service at `/webhooks/confluence` endpoint with signature verification (HMAC-SHA256)
- **Event Queue:** Push events to Redis Streams for async processing; prevents webhook timeout and provides retry capability

Webhook handler pattern:

```

@app.post("/webhooks/confluence")
async def handle_confluence_webhook(
    request: Request,
    background_tasks: BackgroundTasks
):
    # Verify webhook signature
    signature = request.headers.get("X-Confluence-Signature")
    payload = await request.body()
    if not verify_signature(payload, signature, WEBHOOK_SECRET):
        raise HTTPException(401, "Invalid signature")

    event = await request.json()
    event_type = event["webhookEvent"]

    # Queue for async processing
    await redis.xadd(
        "confluence:events",
        {"type": event_type, "data": json.dumps(event)}
    )

    return {"status": "queued"}

```

Incremental Ingestion Worker

- **Content Extraction:** Use `atlassian-python-api` to fetch page content, metadata, and attachments via REST API
- **Change Detection:** Store content hash (SHA-256) in PostgreSQL; skip re-embedding if hash unchanged
- **Chunking Strategy:** Recursive character-based splitting with 512-token chunks, 64-token overlap; preserve Confluence page structure (headers, tables, code blocks) as metadata
- **Embedding:** Batch embed using `text-embedding-3-large` (1536 dimensions); store in Qdrant with metadata payload

Metadata attached to every chunk:

```
{
  "page_id": "123456",
  "page_title": "Authentication Service Spec",
  "space_key": "ENG",
  "space_name": "Engineering",
  "author": "john.doe@company.com",
  "last_modified": "2026-02-26T08:00:00Z",
  "page_url": "https://company.atlassian.net/wiki/spaces/ENG/pages/123456",
  "labels": ["backend", "security", "api"],
  "content_hash": "a3f5b8c...",
  "parent_page_id": "112233",
  "chunk_index": 2,
  "total_chunks": 8
}
```

Deletion Handling

- **Soft Delete:** Mark vectors as `deleted=true` in Qdrant metadata; exclude from search via filter
- **Hard Delete:** Weekly cleanup job purges vectors marked deleted > 30 days
- **Orphan Detection:** Monthly audit detects pages in vector DB but not in Confluence; flags for review

Scheduled Full Reconciliation

- **Weekly Job:** Compare Confluence space inventories against vector DB; detect missed webhook events
- **Re-embedding Trigger:** If Confluence page `last_modified` > vector DB timestamp, trigger re-ingestion
- **Monitoring:** Alert if reconciliation finds > 5% drift between Confluence and vector DB

1.4 Sync Performance Targets

Metric	Target
Webhook → Vector DB latency (P95)	< 30 seconds
New page → searchable	< 1 minute
Updated page → reflected in results	< 2 minutes
Weekly reconciliation runtime	< 2 hours

Metric	Target
Change detection accuracy	> 99.5%

2. Handling Page Attachments

2.1 Challenge

Confluence pages contain:

- **Images/Diagrams:** Architecture diagrams, workflow charts, screenshots, UI mockups
- **Text Documents:** PDFs, Word docs, Excel sheets, CSVs attached for reference

Standard text-only RAG cannot extract information from visual content or document attachments.

2.2 Solution: Multimodal RAG Pipeline

```

flowchart TD
    A[Confluence Page with Attachments] --> B{Attachment Type Detection}
    B -->|Images: PNG, JPG, SVG| C[Image Processing Pipeline]
    B -->|Documents: PDF, DOCX, XLSX| D[Document Extraction Pipeline]
    B -->|Text: MD, TXT| E[Direct Text Extraction]

    C --> C1[Vision LLM - GPT-4o / Gemini Pro Vision]
    C1 --> C2[Generate Structured Description]
    C2 --> C3[Extract Text via OCR - Tesseract]
    C3 --> EMB1[Embed Description + OCR Text]

    D --> D1[Document Parser]
    D1 --> D2[PDF: PyMuPDF / pdfplumber]
    D1 --> D3[DOCX: python-docx]
    D1 --> D4[XLSX: openpyxl]
    D2 & D3 & D4 --> EMB2[Chunk + Embed Document Content]

    E --> EMB3[Chunk + Embed Text]

    EMB1 & EMB2 & EMB3 --> V[(Qdrant Vector Store)]
    C2 --> PG[(PostgreSQL
    Image Descriptions + Links)]

```

2.3 Image and Diagram Handling

Vision-Language Model Integration

- **Model Selection:** Use GPT-4o (multimodal) or Gemini 1.5 Pro for image understanding
- **Structured Extraction Prompt:**

```
VISION_PROMPT = """
You are analyzing an image attached to a Confluence documentation page.
```

Generate a structured description including:

1. Image Type: [diagram, screenshot, chart, photo, mockup, other]
2. Main Components: List all visible elements (boxes, arrows, text labels, icons)
3. Relationships: Describe connections and flow between components
4. Text Content: Extract all visible text (labels, titles, annotations)
5. Purpose: Infer what this image is documenting or explaining

Be precise and technical. This description will be used for semantic search.

Image Context from Page:

```
Title: {page_title}
Section: {section_heading}
Caption: {image_caption}
"""

```

- **Output Example:**

```
{
  "image_type": "architecture_diagram",
  "components": [
    "API Gateway",
    "Authentication Service",
    "User Database",
    "Redis Cache"
  ],
  "relationships": [
    "API Gateway routes requests to Authentication Service",
    "Authentication Service reads from User Database",
    "Authentication Service caches tokens in Redis"
  ],
  "extracted_text": [
    "POST /auth/login",
    "JWT Token",
    "PostgreSQL"
  ],
  "purpose": "System architecture showing authentication flow and data dependencies"
}
```

OCR Fallback for Text-Heavy Images

- **Tesseract OCR:** Extract text from diagrams with labels, flowcharts, annotated screenshots
- **Combine Vision + OCR:** Vision LLM provides semantic understanding; OCR ensures exact keyword matching for technical terms

Image Metadata Enrichment

```
{  
    "attachment_id": "att_789",  
    "attachment_name": "auth-flow-diagram.png",  
    "page_id": "123456",  
    "image_url": "https://company.atlassian.net/.../auth-flow-diagram.png",  
    "vision_description": "Architecture diagram showing...",  
    "ocr_text": "API Gateway, JWT, Redis...",  
    "detected_entities": ["API Gateway", "Redis", "PostgreSQL"],  
    "image_type": "architecture_diagram"  
}
```

2.4 Document Attachment Handling

PDF Documents

- **Text Extraction:** PyMuPDF (fast) or pdfplumber (table-aware) for text content
- **Table Extraction:** Convert tables to Markdown format; preserve structure for better retrieval
- **Metadata:** Extract title, author, creation date from PDF metadata fields

Word Documents (DOCX)

- **Parser:** python-docx extracts text, headings, tables, and embedded images
- **Preserve Structure:** Maintain heading hierarchy as metadata (**h1**, **h2**, **h3**) for context

Excel Sheets (XLSX)

- **Parser:** openpyxl reads cell values, formulas, sheet names
- **Data Structuring:** Convert sheets to Markdown tables or JSON; embed sheet-level summaries
- **Use Case:** Financial reports, data tables, configuration matrices

Chunking Strategy for Attachments

- **Document Chunks:** 512 tokens per chunk with 64-token overlap (same as page content)
- **Image Descriptions:** Embed as single chunk with full vision LLM description + OCR text
- **Cross-Reference:** Link attachment chunks to parent page chunks via parent_page_id metadata

2.5 Attachment Processing Performance

Attachment Type	Processing Time (P95)	Success Rate Target
Images (< 5MB)	< 10 seconds	> 98%
PDFs (< 20 pages)	< 20 seconds	> 95%
DOCX (< 50 pages)	< 15 seconds	> 97%
XLSX (< 10 sheets)	< 10 seconds	> 99%

3. Handling Many Different Spaces

3.1 Challenge

Different Confluence spaces serve different purposes:

- **Engineering Space:** Technical specs, API docs, architecture decisions
- **Product Space:** Feature requirements, user stories, roadmaps
- **HR Space:** Policies, onboarding guides, benefits documentation
- **Sales Space:** Sales playbooks, case studies, pricing sheets

Each space has unique conventions, terminology, document structures, and relevance contexts.

3.2 Solution: Space-Aware RAG with Contextual Metadata

```

flowchart TD
    Q[User Query: 'How do we handle authentication?'] --> R{Space Routing}
    R --> F1[Filter: space_key IN user_accessible_spaces]
    F1 --> H{Hybrid Search}
    H --> V[Vector Search in Qdrant
    with space metadata filters]
    H --> K[Keyword Search in PostgreSQL
    Full-text index]
    V & K --> RR[Cohere Rerank
    with space context]
    RR --> C[Context Assembly]
    C --> SC[Space-Specific Context Injection]
    SC --> LLM[LLM Generation
    with space conventions]
    LLM --> A[Answer with Citations]

```

3.3 Space Metadata Registry

Space Profile Storage (PostgreSQL)

```

CREATE TABLE confluence_spaces (
    space_key VARCHAR(50) PRIMARY KEY,
    space_name VARCHAR(255),
    space_type VARCHAR(50), -- 'engineering', 'product', 'hr', 'sales'
    description TEXT,
    conventions JSONB, -- custom conventions for this space
    primary_topics TEXT[], -- ['authentication', 'deployment', 'ci-cd']
    doc_structure_hints JSONB, -- where specs live, naming patterns
    access_groups TEXT[], -- AD groups with access
    last_updated TIMESTAMP
);

```

Example Space Profile:

```
{
  "space_key": "ENG",
  "space_name": "Engineering",
  "space_type": "engineering",
  "conventions": {
    "spec_location": "All technical specs under 'Specifications' parent page",
    "decision_records": "ADRs stored with 'adr-' prefix in title",
    "code_samples": "Code blocks use Python/JavaScript/Go",
    "terminology": {
      "auth": "authentication service (OAuth2 + JWT)",
      "db": "PostgreSQL primary, Redis cache"
    }
  },
  "primary_topics": ["backend", "infrastructure", "api-design", "security"],
  "doc_structure_hints": {
    "specs_parent_page_id": "112233",
    "runbooks_parent_page_id": "445566"
  }
}
```

3.4 Space-Aware Retrieval Strategy

Query-Time Space Filtering

- **User Context:** Fetch user's accessible spaces from Confluence API or Azure AD group membership
- **Qdrant Filter:** Apply metadata filter `space_key IN ['ENG', 'PROD']` to vector search
- **Boost Relevant Spaces:** If query contains space-specific terminology (e.g., "API spec"), boost `space_type=engineering` results

Qdrant filter example:

```
from qdrant_client.models import Filter, FieldCondition, MatchAny

search_filter = Filter(
    must=[
        FieldCondition(
            key="space_key",
            match=MatchAny(any=user_accessible_spaces)
        ),
        FieldCondition(
            key="deleted",
            match={"value": False}
        )
    ]
)
```

```

results = qdrant_client.search(
    collection_name="confluence_docs",
    query_vector=query_embedding,
    query_filter=search_filter,
    limit=20
)

```

Space-Specific Context Injection

- **Pre-Retrieval:** Inject space conventions into the LLM prompt before generating the answer
- **Example:**

```

SPACE_CONTEXT_TEMPLATE = """
You are answering a question about the {space_name} space in Confluence.

```

Space Context:

- Type: {space_type}
- Primary Topics: {primary_topics}
- Document Conventions: {conventions}

When answering:

- Use terminology specific to this space
- Reference document structure conventions
- Cite page titles and URLs from this space

Retrieved Context:

{retrieved_chunks}

Question: {user_question}

"""

Cross-Space Retrieval

- **Default Behavior:** Search across all accessible spaces; rank by relevance
- **Multi-Space Results:** Group results by space in the UI; show top 3 results per space
- **Disambiguation:** If query matches multiple spaces (e.g., "authentication" in Engineering + Product), surface both and ask user to clarify

3.5 Space Management Operations

Operation	Implementation	Frequency
Space Profile Creation	Admin creates profile via UI; stores in PostgreSQL	On new space creation
Convention Updates	Admin edits conventions via UI; triggers re-ranking calibration	As needed
Access Sync	Fetch user → space mappings from Confluence API	Every 6 hours

Operation	Implementation	Frequency
Space-Level Analytics	Track query distribution, answer quality per space	Daily aggregation

4. Testing Answer Quality

4.1 Challenge

LLM-generated answers can suffer from:

- **Hallucinations:** Fabricating information not present in retrieved context
- **Irrelevance:** Answering a different question than asked
- **Incompleteness:** Missing key details from retrieved documents
- **Staleness:** Using outdated information despite having current data

Continuous evaluation is critical for maintaining user trust.

4.2 Evaluation Framework

```
flowchart TD
    A[RAG System Generates Answer] --> B[Automated Evaluation Layer]
    A --> C[Human Feedback Loop]

    B --> B1[LLM-as-Judge - GPT-4o]
    B1 --> M1[Faithfulness Score]
    B1 --> M2[Answer Relevance Score]
    B1 --> M3[Contextual Precision]
    B1 --> M4[Completeness Score]

    C --> C1[User Rating:  / ]
    C --> C2[Expert Review Queue]

    M1 & M2 & M3 & M4 & C1 & C2 --> D[(Evaluation DB PostgreSQL)]
    D --> E[Weekly Quality Report]
    E --> F{Score Regression?}
    F -->|Yes| G[Trigger Improvement Loop]
    F -->|No| H[Continue Monitoring]

    G --> I[Prompt Engineering]
    G --> J[Retrieval Tuning]
    G --> K[Fine-tune Reranker]
```

4.3 Automated Evaluation Metrics

Using RAGAS Framework

```

from ragas import evaluate
from ragas.metrics import (
    faithfulness, # Is answer grounded in context?
    answer_relevancy, # Is answer relevant to question?
    context_precision, # Are retrieved docs relevant?
    context_recall, # Did we retrieve all needed info?
    answer_correctness # Compare against ground truth
)

# Evaluation dataset structure
eval_data = {
    "question": ["How do we authenticate API requests?"],
    "answer": ["API requests use JWT tokens issued by..."],
    "contexts": [[["Context chunk 1..."], ["Context chunk 2..."]]],
    "ground_truth": ["Authentication is handled via OAuth2..."]
}

# Run evaluation
result = evaluate(
    eval_data,
    metrics=[# Is answer grounded in context?
        faithfulness,
        answer_relevancy,
        context_precision,
        context_recall,
        answer_correctness
    ]
)

# Result: {'faithfulness': 0.89, 'answer_relevancy': 0.92, ...}

```

LLM-as-Judge Scoring (vRAG-Eval)

5-Point Grading Scale:

Score	Meaning	Criteria
5	Excellent	Correct + complete + adds helpful context
4	Correct	Answers question accurately from retrieved docs
3	Borderline	Partially correct or contains minor errors
2	Incorrect	Wrong information or misinterprets question
1	Hallucination	Fabricates information not in context

Judge Prompt Template:

```

JUDGE_PROMPT = """
You are evaluating the quality of an answer generated by a RAG system.

```

Question: {question}

Retrieved Context:
{context_chunks}

Generated Answer:
{generated_answer}

Ground Truth (if available):
{ground_truth}

Rate the answer on a scale of 1-5:

- 1 = Hallucination (fabricated info)
- 2 = Incorrect (wrong information)
- 3 = Borderline (partially correct, minor errors)
- 4 = Correct (accurate from context)
- 5 = Excellent (correct + complete + helpful)

Provide:

1. Score (1-5)
2. Justification (2-3 sentences)
3. Identified Issues (if score < 4)

Output as JSON:

```
{  
  "score": 4,  
  "justification": "Answer correctly uses information from context...",  
  "issues": []  
}  
"""
```

4.4 Human Feedback Loop

Inline User Feedback

- **UI Component:** After every answer, show / buttons + optional comment field
- **Tracking:** Store (question, answer, rating, comment, user_id, timestamp) in PostgreSQL
- **Weekly Aggregation:** Calculate acceptance rate, negative feedback trends per space

Expert Review Queue

- **Trigger Conditions:** Auto-queue answers with:
 - LLM-as-judge score < 3.5
 - User downvote + comment
 - High-stakes queries (identified by keywords: "policy", "legal", "security")
- **Review Process:** Subject matter experts (SMEs) grade answers on same 1-5 scale; provide corrected answer if needed

- **Golden Dataset:** Add expert-approved (question, context, answer) triplets to regression test suite

4.5 Continuous Improvement Loop

Weekly Quality Dashboard

```

flowchart LR
    D[Evaluation DB] --> A[Aggregate Metrics]
    A --> M1[Avg Faithfulness: 0.88 ↓]
    A --> M2[User Acceptance: 82% ↑]
    A --> M3[Expert Review: 15 pending]
    A --> M4[P95 Latency: 6.2s]

    M1 & M2 & M3 & M4 --> R{Threshold Breach?}
    R -->|Yes| T[Trigger Investigation]
    R -->|No| OK[Continue Monitoring]

    T --> I1[Analyze Low-Scoring Queries]
    I1 --> I2[Identify Patterns]
    I2 --> I3[Root Cause]

    I3 --> F1[Prompt Engineering]
    I3 --> F2[Retrieval Tuning]
    I3 --> F3[Reranker Fine-tuning]
    I3 --> F4[Space Profile Update]

    F1 & F2 & F3 & F4 --> V[Validate in Staging]
    V --> P[Promote to Production]

```

Threshold-Based Alerts

Metric	Alert Threshold	Action
Faithfulness score (weekly avg)	< 0.80	Review prompt engineering + context window size
User acceptance rate	< 75%	Analyze negative feedback patterns; SME review
Hallucination rate	> 10%	Strengthen guardrails; reduce temperature
Answer relevance	< 0.85	Tune retrieval filters + reranker weights
Expert review backlog	> 25 items	Allocate additional SME time

4.6 A/B Testing for Improvements

- **Prompt Versioning:** Test new system prompts on 20% of traffic in production
- **Retrieval Strategy:** A/B test hybrid search weights (vector vs. keyword ratio)
- **Reranker Models:** Compare Cohere vs. bge-reranker on sampled queries
- **Promotion Gate:** Promote only if new version improves faithfulness + relevance by $\geq 3\%$

4.7 Regression Test Suite

- **Golden Dataset:** Curate 100 question-answer pairs validated by experts
- **Coverage:** Include queries spanning all spaces, attachment types, and complexity levels
- **CI Gate:** Run full evaluation on every deployment; block if scores drop > 5%
- **Maintenance:** Add 5-10 new examples monthly from expert review queue

4.8 Quality Targets

Metric	Target	Measurement
Faithfulness	≥ 0.85	RAGAS automated eval
Answer Relevance	≥ 0.88	RAGAS automated eval
User Acceptance Rate	≥ 80%	👍 / (👍 + 🤦)
Hallucination Rate	< 5%	LLM-as-judge score = 1
Expert Approval Rate	≥ 90%	SME review queue
Context Precision	≥ 0.80	RAGAS retrieval eval
P95 Answer Latency	< 8 seconds	Prometheus metrics

5. System Architecture Overview

```

flowchart TD
    subgraph Users ["👤 End Users"]
        U1[Web UI - React]
        U2[Slack Bot]
        U3[Teams Bot]
        U4[REST API Clients]
    end

    subgraph API ["⚡ API Layer - FastAPI"]
        A1[Search Endpoint]
        A2[Feedback Endpoint]
        A3[Admin Dashboard]
        A4[Webhook Receiver]
    end

    subgraph Orchestration ["⚙️ RAG Orchestration"]
        O1[Query Rewriter]
        O2[Hybrid Retriever]
        O3[Cohere Reranker]
        O4[Context Assembler]
        O5[LLM Gateway - LiteLLM]
        O6[Guardrails Layer]
    end

    subgraph Storage ["📁 Data Layer"]
        S1[(Qdrant Vector DB)]
        S2[(PostgreSQL)]
    end

```

```

        Metadata + Feedback)]
S3[(Redis
Cache + Queue)]
end

subgraph Ingestion["📄 Ingestion Pipeline"]
I1[Confluence Sync Worker]
I2[Image Processor]
I3[Document Processor]
I4[Embedder]
end

subgraph Eval["📊 Evaluation"]
E1[LLM-as-Judge]
E2[Expert Review Queue]
E3[Quality Dashboard]
end

Users --> API
API --> Orchestration
Orchestration --> Storage
A4 --> I1
I1 --> I2 & I3 --> I4 --> Storage
Orchestration --> E1
API --> E2
E1 & E2 --> E3

```

6. Implementation Roadmap

Phase	Duration	Deliverables
Phase 1: Foundation	3 weeks	Confluence API integration, webhook setup, basic sync pipeline, PostgreSQL schema, Qdrant setup
Phase 2: Core RAG	3 weeks	Hybrid retrieval, reranking, LLM integration (LiteLLM), basic Q&A endpoint, web UI prototype
Phase 3: Multimodal	2 weeks	Vision LLM integration, image description pipeline, document parser (PDF/DOCX/XLSX), attachment handling
Phase 4: Space-Aware	2 weeks	Space metadata registry, space-specific context injection, access control integration, cross-space search
Phase 5: Evaluation	2 weeks	RAGAS integration, LLM-as-judge pipeline, user feedback UI, expert review queue, quality dashboard
Phase 6: Hardening	2 weeks	Load testing, security audit, monitoring (Prometheus + Grafana), guardrails, error handling
Phase 7: Pilot	2 weeks	20-user pilot, feedback collection, iteration, final tuning

Phase	Duration	Deliverables
Total	16 weeks	Production-ready Confluence Q&A system

7. Success Metrics

Category	Metric	Target
Adoption	Monthly active users (MAU)	\geq 60% of organization
Engagement	Queries per user per week	\geq 5
Quality	User acceptance rate ($\frac{\text{✓}}{\text{✗}}$ / total)	\geq 80%
Quality	Expert approval rate	\geq 90%
Quality	Faithfulness score (RAGAS)	\geq 0.85
Performance	P95 query latency	< 8 seconds
Performance	Sync lag (Confluence → searchable)	< 2 minutes
Reliability	System uptime	\geq 99.5%
Cost	Cost per query	< \$0.03

8. Risk Mitigation

Risk	Impact	Mitigation
Confluence API rate limits	High	Implement exponential backoff, request batching, cache frequent queries
Vision LLM cost	Medium	Process only images > 50KB, cache descriptions, use cheaper models for low-priority images
Hallucination incidents	High	Strengthen guardrails, mandatory citations, flag low-confidence answers for review
Space access violations	Critical	Enforce Confluence permissions at query time, audit access logs weekly
Vector DB performance degradation	Medium	Horizontal scaling, read replicas, query optimization, index tuning
Stale content in results	Medium	Real-time sync with SLA monitoring, weekly reconciliation, content freshness indicator in UI

9. Technology Stack Summary

Layer	Technology
-------	------------

Layer	Technology
Orchestration	LangChain, LangGraph (optional for complex flows)
Model Gateway	LiteLLM
Embeddings	OpenAI text-embedding-3-large
Vision LLM	GPT-4o or Gemini 1.5 Pro
Vector Store	Qdrant
Metadata DB	PostgreSQL with pgvector extension
Cache & Queue	Redis (Streams + Cache)
Reranker	Cohere Rerank API
Document Parsing	PyMuPDF, python-docx, openpyxl, Tesseract OCR
API Framework	FastAPI + Unicorn
Evaluation	RAGAS, DeepEval, custom LLM-as-judge
Monitoring	Prometheus, Grafana, Langfuse
Integration	atlassian-python-api for Confluence API
Containerization	Docker, Kubernetes (AKS)

10. Next Steps

- Week 1:** Stakeholder alignment, finalize Confluence access permissions, provision infrastructure (Qdrant, PostgreSQL, Redis)
- Week 2:** Set up webhook integration, build sync pipeline MVP, test incremental ingestion
- Week 3:** Implement hybrid retrieval + basic Q&A endpoint, deploy staging environment
- Week 4+:** Follow phased roadmap; weekly demos to stakeholders; iterate based on feedback

Document Owner: AI Engineering Team

Last Updated: February 26, 2026

Status: Architecture Approved — Awaiting Implementation Kickoff